



How to Meta-Sumo

Using Logs for Agile Monitoring of
Production Services

Christian Beedgen

christian@sumologic.com

[@raychaser](#)

Who Am I?

- Co-Founder & CTO, Sumo Logic since 2010
 - Cloud-based log management and analytics
 - Applications, operations, security
- Server guy, Chief Architect, ArcSight, 2001 – 2009
 - Major SIEM player in the enterprise space
 - Log management for security and compliance



NPR Business

@nprbusiness



Follow



Hewlett-Packard to lay off 27,000 workers,
8 percent of its work force



Reply



Retweet



Favorite



NPR Business

@nprbusiness



Follow



Hewlett-Packard to lay off 27,000 workers,
8 percent of its work force



Reply



Retweet



Favorite



NPR Business

@nprbusiness



Follow



HP sees restructuring saving up to \$3.5
billion; to invest in the "cloud, big data and
security."

Logs & Logging

What are Logs?

- The murmurings and whispers of your infrastructure
 - Devices & Services (ex. Security, Email, Authentication)
 - Applications (your code, or 3rd party applications)
- Written to disk by applications
 - Used during development by developers
 - And then often ignored in production
- Yes, Logs are Big Data
 - There's a lot of Logs (→ Volume)
 - A large Variety of formats, plus it's real-time (→ Velocity)
- Free-form text, usually one message per line
 - Scrolls by in your terminal
 - Makes you feel like you're in the matrix

Application Logging

- **Just do it***
 - Use what's common in your language (Log4J, ...)
 - If in doubt, log it - you might need it later, and disk is cheap
- **The obvious stuff**
 - Always log a precise timestamp
 - Add a log level for filtering
- **Building a distributed system?**
 - Always add the process/service/module name
 - Add the host ID if you can
- **Log the context**
 - When processing a request, remember who it is from
 - Within the scope of the request, always log the context
- **Bring it all together in a single place**
 - Open Source options
 - Commercial options

Just Do It*

- Actually, please do think first
 - Do not log *passwords*
 - Do not log *credit card numbers*
 - Do not log *anything that's PII*

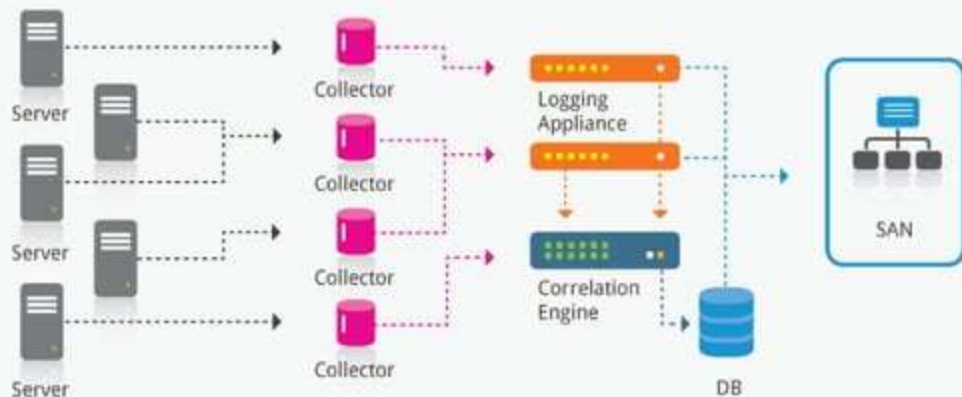
Or Else...



<http://www.slideshare.net/christoferhoff/shit-my-cloud-evangelist-says-just-not-to-my-cso>

Log Management


Traditional On-premise Deployment



Log Management

Sumo Logic Cloud-based Deployment


Customer A
Data Center



Customer B
Data Center



Customer C
Data Center



Sumo Logs


This is the Meta part

- Our system to collect and analyze Logs is actually a distributed, cloud-based infrastructure that itself generates a ton of logs
- We are running a second, smaller instance that receives all the logs from the **Prod** system – we call this our **Shadow** system
- A majority of the management and monitoring of the **Prod** system is done via the **Shadow** system, which we are madly in love with

My team



My team




No, really. And we also have a Panda.

Example

```
2012-05-22 18:47:26,807 -0700 INFO [hostId=long-frontend-1] [module=RECEIVER]
[logger=scala.receiver.MessageBlocker] [thread=MTP-MessagePilePipeline-3]
[auth=Collector:prod-cass-raw-8:000000000000483D:0000000000000005:false]
[remote_ip=184.73.74.54] [web_session=MepMG8CS...] Pile for customer:
'0000000000000005', ID: '800000006407637B', block: '80000000004C9A11', msg
count: '1', size: '264', collector: '000000000000483D'
```


Example

```
2012-05-22 18:47:26,807 -0700 [long-frontend-1] [module=RECEIVER]
[logger=scala.receiver.MessageBuffer] [thread=MTP-MessagePilePipeline-3]
[auth=Collector:prod-cass-raw-8:000000000000483D:0000000000000005:false]
[remote_ip=184.73.74.54] [web_session=MepMG8CS...] Pile for customer:
'0000000000000005', ID: '800000006407637B', block: '80000000004C9A11', msg
count: '1', size: '264', collector: '000000000000483D'
```



- Timestamp with time zone!

Example

```
2012-05-22 18:47:26,807 -0700 INFO [frontend-1] [module=RECEIVER]
[logger=scala.receiver.MessageBlocker, [thread=etf-MessagePipeline-3]
[auth=Collector:prod-cass-raw-8:0000000000000483D:0000000000000005:false]
[remote_ip=184.73.74.54] [web_session=MepMG8CS...] File for customer:
'00000000000000005', ID: '800000006407637B', block: '80000000004C9A11', msg
count: '1', size: '264', collector: '0000000000000483D'
```



- Timestamp with time zone!
- Log level

Example

```
2012-05-22 18:47:26, [hostId=long-frontend-1] [module=RECEIVER]  
[logger=scala.receiver.messagesucker] [thread=MTP-MessagePilePipeline-3]  
[auth=Collector:prod-cass-raw-8:000000000000483D:0000000000000005:false]  
[remote_ip=184.73.74.54] [web_session=MepMG8CS...] File for customer:  
'000000000000000005', ID: '800000006407637B', block: '80000000004C9A11', msg  
count: '1', size: '264', collector: '000000000000483D'
```



- Timestamp with time zone!
- Log level
- Host ID & module name (process/service)

Example


```
2012-05-22 18:47:26,807 -0700 INFO [hostId=long-frontend-1] [module=RECEIVER]
[logger=scala.receiver.MessageBlocker] [messagePipeline=3]
[auth=Collector:prod-cass-raw-8:0000000000000483D:0000000000000005:false]
[remote_ip=184.73.74.54] [web_session=MepMG8CS...] File for customer:
'000000000000000005', ID: '800000006407637B', block: '80000000004C9A11', msg
count: '1', size: '264', collector: '0000000000000483D'
```



- Timestamp with time zone!
- Log level
- Host ID & module name (process/service)
- Code location or class

Example

```
2012-05-22 18:47:26,807 -0700 INFO [hostId=long-frontend-1] [module=RECEIVER]
[logger=scala.receiver.MessageBlocker] [thread=MTP-MessagePilePipeline-3]
[auth=Collector:prod-cass-raw-8:000000000000483D:0000000000000005:false]
[remote_ip=184.73.74.54] [web_session=MepMG8CS...] Pile for customer:
'0000000000000005', ID: '800000006407637B', block: '80000000004C9A11', msg
count: '1', size: '264', collector: '000000000000483D'
```



- Timestamp with time zone!
- Log level
- Host ID & module name (process/service)
- Code location or class
- Authentication context

Example

```
2012-05-22 18:47:26,807 -0700 INFO [hostId=long-frontend-1] [module=RECEIVER]
[logger=scala.receiver.MessageBlocker] [thread=MTP-MessagePilePipeline-3]
[auth=Collector:prod-cass-raw-8:000000000000483D:0000000000000005:false]
[remote_ip=184.73.74.54] [web_session=MepMG8CS...] File for customer:
'0000000000000005', ID: '800000006407637B', block: '80000000004C9A11', msg
count: '1', size: '264', collector: '000000000000483D'
```

- Timestamp with time zone!
- Log level
- Host ID & module name (process/service)
- Code location or class
- Authentication context
- Key-value pairs



What about structure?

- Should you rather use JSON, say?
 - Hey, it's really up to you – machines love it
 - But do humans read the logs as well?
- Modern tools extract structure if there is structure
 - Even if the structure doesn't conform to a protocol
 - So you don't strictly have to use a structured format

Sumo Pet Tricks

Sumo Pet Trick #1 – Basic Troubleshooting

- All errors in the application have an error instance ID



- UI always tries very hard to display the ID
 - User can easily copy/paste the error ID into a ticket

Sumo Pet Trick #1 – Basic Troubleshooting

- Find AP2V3-HZICT-BP6UO in the logs!

```

2012-05-08 11:00:41.974 -0700 ERROR [hostid-prod-frontend-3] [module=SERVICE] [logger=service.util.exception.ExceptionHandlingInterceptor] [thread=btpool10
[auth-user=: : false] [remote-ip=: ] [web_session=thisfzle...] Error ID "APZV3-H2ICT-BPMUO" while
ReflectiveMethodInvocation: public abstract com.sumologic.service.endpoint.user.v1.model.AddUserResult
com.sumologic.service.endpoint.user.v1.api.UserManagementService.addUser(java.lang.String,com.sumologic.service.endpoint.user.v1.model.User) set is of cl
[com.sumologic.service.endpoint.user.v1.impl.UserManagementServiceImpl]'
java.lang.RuntimeException: java.util.IllegalFormatConversionException: c != java.lang.String
    at com.sumologic.service.endpoint.util.v1.model.ServiceUtils$.safe(ServiceUtils.scala:37)
    at com.sumologic.service.endpoint.user.v1.impl.UserManagementServiceDelegate.$anonfun$addUser(UserManagementServiceDelegate.scala:29)
    at com.sumologic.service.endpoint.user.v1.impl.UserManagementServiceDelegate.$anonfun$getOrCreateObject(UserManagementServiceDelegate.scala:68)
    at com.sumologic.service.endpoint.user.v1.impl.UserManagementServiceDelegate.$anonfun$userExistsInCurrentOrg$(UserManagementServiceDelegate.scala:143)
    at com.sumologic.service.endpoint.user.v1.impl.UserManagementServiceDelegate.$anonfun$findUserExistsInCurrentOrg$(UserManagementServiceDelegate.scala:143)
    at com.sumologic.config.graphORM.Transaction$.nonFunTransaction$lzycompute(Transaction.scala:680)
    at com.sumologic.config.graphORM.Transaction$.Metrics$$anonfun$withGlobalTimer$(Transaction.scala:143)
    at com.yammer.metrics.Timer.time(Timer.scala:20)
    at com.sumologic.util.scala.Metrics$class.withGlohalTimer(Metrics.scala:136)
    at com.sumologic.config.graphORM.Transaction$.withGlobalTimer(TransactionalTransaction.scala:29)
    at com.sumologic.util.scala.Metrics$$anonfun$withOrganizationTimer$(Metrics.scala:112)
    at com.yammer.metrics.Timer.time(Timer.scala:20)
    at com.sumologic.util.scala.Metrics$class.withOrganizationTimer(Metrics.scala:111)
    at com.sumologic.config.graphORM.Transaction$.withOrganizationTimer(TransactionalTransaction.scala:29)
    at com.sumologic.util.scala.Metrics$class.time(Metrics.scala:29)
    at com.sumologic.config.graphORM.Transaction$.time(Transaction.scala:29)
    at com.sumologic.config.graphORM.Transaction$.transaction(transaction.scala:657)
    at com.sumologic.service.endpoint.user.v1.impl.UserManagementServiceDelegate.userExistsInCurrentOrg(UserManagementServiceDelegate.scala:150)
    at com.sumologic.service.endpoint.user.v1.impl.UserManagementServiceDelegate.addUser(UserManagementServiceDelegate.scala:205)
    at com.sumologic.service.endpoint.user.v1.impl.UserManagementServiceImpl addUser(UserManagementServiceImpl.java:49)
    at sun.reflect.NativeMethodAccessorImpl.invoke(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:39)
    at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:25)
    at java.lang.reflect.Method.invoke(Method.java:597)

```


Sumo Pet Trick #2 – Everything breaks

- Distributed systems → lots of different components
 - There will be fail
 - The morning coffee routine...
- There is beauty in numbers
 - Pull out all the logs that contain the term “error” (or “fail”, or “exception”, or all of the above)
 - Then count by process/service/module, or host ID
- Even the most basic aggregation will create actionable insight

Sumo Pet Trick #2 – Everything breaks

```
error | parse "hostId=*" as hostId | count as error_count by hostId | sort error_count
```

#	hostId	error_count
1	long-config-1	1,633
2	long-ftsearch-2	298
3	long-ftsearch-3	236
4	long-ftsearch-1	201
5	long-cass-raw-3	64
6	long-cass-raw-1	57
7	long-cass-raw-2	47
8	long-frontend-2	27
9	long-frontend-1	22
10	long-cass-meta-3	18
11	long-cass-meta-1	16
12	long-cass-meta-2	17
13	long-bill-1	1



#	module	error_count
1	CONFIG	1,645
2	SEARCH	425
3	STREAM	315
4	RAW	170
5	META	54
6	SERVICE	49
7	BILL	1



```
error | parse "module=*" as module | count as error_count by module | sort by error_count
```

Sumo Pet Trick #2 – Everything breaks

```
error | parse "hostId=*" as hostId | count as error_count by hostId | sort error_count
```

#	hostId	error_count
1	long-config-1	1,633
2	long-ftsearch-2	298
3	long-ftsearch-3	236
4	long-ftsearch-1	201
5	long-cass-raw-3	64
6	long-cass-raw-1	57
7	long-cass-raw-2	47
8	long-frontend-2	27
9	long-frontend-1	22
10	long-cass-meta-3	18
11	long-cass-meta-1	18
12	long-cass-meta-2	17
13	long-bill-1	1

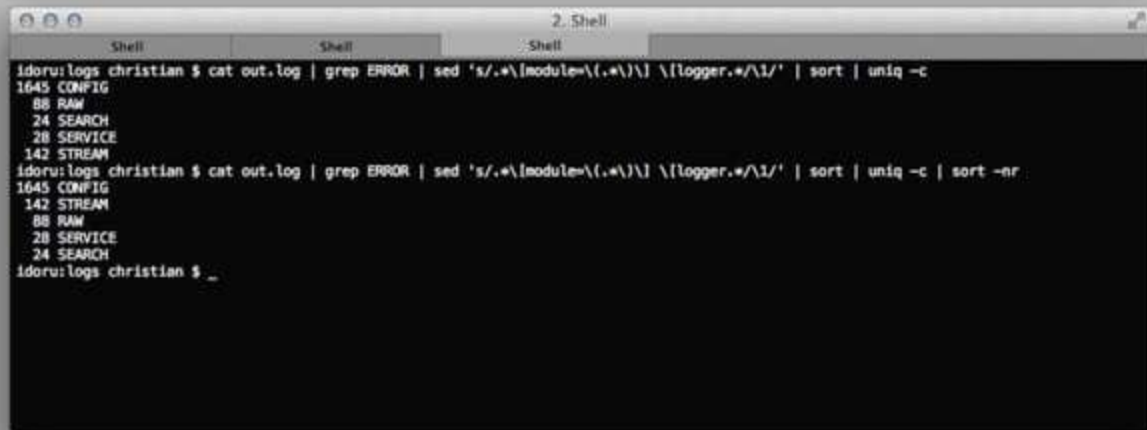
Schedule
this!

#	module	error_count
1	CONFIG	1,645
2	SEARCH	425
3	STREAM	315
4	RAW	170
5	META	54
6	SERVICE	49
7	BILL	1

```
error | parse "module=*" as module | count as error_count by module | sort by error_count
```

Sumo Pet Trick #2 – Everything breaks

- Old school can rock this too!



```
2. Shell
Shell Shell Shell
idoru:logs christian $ cat out.log | grep ERROR | sed 's/.*[module=\\(.\\)\\] \\logger.*/\\1/' | sort | uniq -c
1645 CONFIG
88 RAW
24 SEARCH
28 SERVICE
142 STREAM
idoru:logs christian $ cat out.log | grep ERROR | sed 's/.*[module=\\(.\\)\\] \\logger.*/\\1/' | sort | uniq -c | sort -nr
1645 CONFIG
142 STREAM
88 RAW
28 SERVICE
24 SEARCH
idoru:logs christian $ _
```

Sumo Pet Trick #3 – API Choke Points

- What is my API doing?
 - Which APIs are being called?
 - Average execution times per API
 - Who's calling?

```
5:42,279 -0700 INFO [hostId=long-frontend-2] [module=SERVICE] [logger=service.util.interceptor.Loggi  
-14] [auth=User:peter@demo.com:000000000000195C:0000000000000005:false] [remote_ip=98.248.40.103] [we  
flectiveMethodInvocation: public abstract com.sumologic.service.endpoint.search.v2.model.GetMessagesR  
service.endpoint.search.v2.api.SearchQueryService.getMessage(java.lang.String,java.lang.String,java.l  
ass [com.sumologic.service.endpoint.search.v2.impl.SearchQueryServiceImpl]', arguments:  
37AE93195A4,raw,1211,15,true', elapsed: '73'
```

Sumo Pet Trick #3 – API Choke Points

- Number of calls by API

#	api	_count
1	search.v1.api.SearchService.getStatus	36,799
2	collector.v1.api.CollectorService.getHistogramForCollector	1,110
3	search.v2.api.SearchQueryService.getSearchQueryStatus	1,089
4	search.v1.api.SearchService.startQuery	582
5	auth.v1.api.AuthenticationService.logOut	581
6	search.v1.api.SearchService.getMessages	300
7	search.v1.api.SearchService.cancelQuery	300
8	collector.v1.api.CollectorService.getCollectorsWithBlades	216
9	account.v1.api.UserPreferencesService.updateInitialTab	149
10	search.v2.api.SearchQueryService.createSearchQuery	112
11	account.v1.api.UserPreferencesService.getUserPreferences	102
12	auth.v1.api.AuthenticationService.getCurrentUserInfo	95
13	account.v1.api.AccountService.getSubscriptionDetails	93
14	search.v2.api.SearchQueryService.deleteSearchQuery	91

Sumo Pet Trick #3 – API Choke Points

- Average execution time by API

#	api	avg_ms
1	collector.v1.api.CollectorService.getHistogramForCustomer	1,405
2	collector.v1.api.CollectorService.getHistogramForCollector	1,199
3	search.v1.api.SavedSearchService.getSearches	590
4	account.v1.api.AccountService.getSubscriptionAndBillingUsages	547
5	user.v1.api.UserManagementService.listUsers	319
6	account.v1.api.AccountService.getSubscriptionDetails	247
7	collector.v1.api.CollectorService.collectorDownloadOfferings	173
8	search.v2.api.SearchQueryService.createSearchQuery	170
9	search.v1.api.SearchService.getMessages	166
10	search.v1.api.SearchService.startQuery	136
11	collector.v1.api.CollectorService.getCollectorsWithBlades	110
12	auth.v1.api.AuthenticationService.getCurrentUserInfo	91
13	search.v2.api.SearchQueryService.parseSearchQuery	91
14	collector.v1.api.CollectorService.createWildcardLocalFileTailBlade	78

Sumo Pet Trick #3 – API Choke Points

- Number of API calls by user

#	user	api	_count
1	test@demo.com	search.v1.api.SearchService.getStatus	37,956
2	mandy.liu@demo.com	collector.v1.api.CollectorService.getHistogramForCollector	1,044
3	rishi@sumologic.com	search.v2.api.SearchQueryService.getSearchQueryStatus	694
4	test@demo.com	search.v1.api.SearchService.startQuery	594
5	test@demo.com	auth.v1.api.AuthenticationService.logout	593
6	test@demo.com	search.v1.api.SearchService.cancelQuery	305
7	test@demo.com	search.v1.api.SearchService.getMessages	305
8	james+example@sumologic.com	search.v2.api.SearchQueryService.getSearchQueryStatus	283
9	mandy.liu@demo.com	collector.v1.api.CollectorService.getCollectorsWithBlades	184
10	mandy.liu@demo.com	account.v1.api.UserPreferencesService.updateInitialTab	91
11	peter@demo.com	search.v2.api.SearchQueryService.getSearchQueryStatus	87
12	mandy.liu@demo.com	account.v1.api.UserPreferencesService.getUserPreferences	85
13	rishi@sumologic.com	search.v2.api.SearchQueryService.createSearchQuery	84
14	mandy.liu@demo.com	auth.v1.api.AuthenticationService.getCurrentUserInfo	81

Sumo Pet Trick #4 – Which searches are running

- Log a session ID on start and stop
 - When there's long-running operations
 - More than one node participates in generating the result
- For each session in the time range, did you get both the start and the stop message?

```
"module=SERVICE" SearchServiceImpl ("Started query" or "canceling query")
| parse "auth=User:*" as user
| parse regex "for session.*" (?<session>.*?) ""
| parse "query: '" as query nodrop
| count, first(query) by session, user
| where _count != 2
```

Sumo Pet Trick #4 – Which searches are running

#	session	user	_count	_first
1	7DF97970DEA03645	test@demo.com	1	etime AND lconnect AND ldisconnect IQA_EXTENDED_SEARCH_LOG
2	0277A1F1570F4E51	panda@demo.com	1	93025E32B0315B6A _sourcecategory=stream "Raw finished processing request" parse "messageCount:" as msg avg(msg), sum(msg), max(msg)
3	5033375B6741B2B7	test@demo.com	1	(_sourceHost=prod-fssearch-4 RawResolverDelegate) AND Received IQA_EXTENDED_SEARCH_LOG parse "sessionId:" as sessionId parse "size:" as size sum(size) by sessionId
4	111095253268E486	peter@demo.com	1	5F20D09072682F69 concierge executing query parse "Executing query for stream session:"
5	025AFAF5ABD8AB40	peter@demo.com	1	93025E32B0315B6A concierge executing query parse "Executing query for stream session:"
6	6FD7842D42552D44	test@demo.com	1	"Summarize operator" IQA_EXTENDED_SEARCH_LOG parse "(maxDepth=", totalNodes=", maxClusterId=") as maxDepth,totalNodes,maxClusterId, "searchExpression=" as searchExpression toLong(maxClusterId) sort by maxClusterId
7	3FAB78A0D505D64	test@demo.com	1	error "Error while removing index" INFO "session has already been cancelled" ICassandraMessagesRetriever "Could not remove pending" IQA_EXTENDED_SEARCH_LOG
8	1A9EB9314A849B6D	rishi@sumologic.com	1	
9	279FFAE94BC4B105	peter@demo.com	1	6156D23ED7ED8BCD SearchServiceImpl creating query parse "session:"
10	0378709A8E1791CD	rishi@sumologic.com	1	_sourceHost=px extract "(?<PXX-id-id(8))" count by msg_code
11	3C500CF0915B8EB2	test@demo.com	1	error or fail * IQA_EXTENDED_SEARCH_LOG summarize
12	676A1FF7FF22A893	peter@demo.com	1	(2E2B0A1F5D40347D or 0B2C07B49451BA47 or 96A207EA7D7E606D or 93025E32B0315B6A or 4AAA86552FC4334C)
13	6D86CC5E0D5EE63F	test@demo.com	1	"logger=org.quartz.core.ErrorLogger" "failed to execute task" IQA_EXTENDED_SEARCH_LOG parse regex "(?<Caused by: java(?.*)?)\n" nodrop parse regex "(?<Caused by: java(?.*)?)\n" sessionId: "(?<n>)" nodrop parse regex "(?<Caused by: java.util.concurrent.ExecutionException: java.lang.RuntimeException: Unable to run query "(?<?.*)?"\n)"\n" nodrop count by cause sort _count

Sumo Pet Trick #4 – Which searches are running

#	session	user	_count	_first
1	7DF97970EA03645	test@demo.com	1	etime AND lconnect AND ldisconnect IQA_EXTENDED_SEARCH_LOG
2	0277A1F1570F4E51	panta@demo.com	1	93025E32B0315B6A _sourcecategory=stream "Raw finished processing request" parse "messageCount:" as msg avg(msg), sum(msg), max(msg)
3	503337586741B2B7	test@demo.com	1	(_sourceHost=prod-fssearch-4 RawResolverDelegate) AND Received IQA_EXTENDED_SEARCH_LOG parse "sessionId:" as sessionId parse "size:" as size sum(size) by sessionId
4	11109525328E486	peter@demo.com	1	5F20D09072682F89 concierge executing query parse "Executing query for stream session:"
5	025AFAF5ABD9AB40	peter@demo.com	1	93025E32B0315B6A concierge executing query parse "Executing query for stream session:"
6	6FD7842042552D44	test@demo.com	1	"Summarize operator" IQA_EXTENDED_SEARCH_LOG parse "[maxDepth=, totalNodes=, maxClusterId=]" as maxDepth,totalNodes,maxClusterId,"searchExpression=" as searchExpression toLong(maxClusterId) sort by maxClusterId
7	3FAB786A2D505064	test@demo.com	1	error !Error while removing index" !INFO !"session has already been cancelled" !CassandraMessagesRetriever !"Could not remove pending" IQA_EXTENDED_SEARCH_LOG
8	1A9EB9314A849B6D	nishi@sumologic.com	1	
9	279FFAE94BC4B105	peter@demo.com	1	8156D23ED7ED8BCD SearchServiceImpl creating query parse "session:"
10	0378709A8E1791CD	nishi@sumologic.com	1	_sourceHost=px extract "[?%PX-d-id(6)]" count by msg_code
11	3C500CF091588EB2	test@demo.com	1	error or fail * IQA_EXTENDED_SEARCH_LOG summarize
12	676A1FF7FF22A893	peter@demo.com	1	(2E2B0A1F5D40347D or 0B2CD7849451BA47 or 96A207EA7D7E606D or 93025E32B0315B6A or 4AA88552FC4334C)
13	6D85CC5E0D5EE63F	test@demo.com	1	"logger=org.quartz.core.ErrorLogger" failed to execute task" IQA_EXTENDED_SEARCH_LOG parse regex ".*inCaused by: java(?.*)n" nodrop parse regex ".*inCaused by: java(?.*) sessionid: .?n" nodrop parse regex ".*inCaused by: java.util.concurrent.ExecutionException: java.lang.RuntimeException: Unable to run query .* (?.*)n" nodrop count by cause sort _count

Thank you!

<https://www.sumologic.com/free-trial/>

Christian Beedgen
christian@sumologic.com
@raychaser

